

Ensemble Learning

Instructor: Alan Ritter

Slides Adapted from Carlos Guestrin and Luke Zettlemoyer

Voting (Ensemble Methods)


- Instead of learning a single classifier, **learn many weak classifiers** that are **good at different parts of the data**
- **Output class:** (Weighted) vote of each classifier
 - Classifiers that are most “sure” will vote with more conviction
 - Classifiers will be most “sure” about a particular part of the space
 - On average, do better than single classifier!
- **But how???**
 - force classifiers to learn about different parts of the input space? different subsets of the data?
 - weigh the votes of different classifiers?

BAGGing = Bootstrap AGGregation

(Breiman, 1996)


- for $i = 1, 2, \dots, K$:
 - $T_i \leftarrow$ randomly select M training instances with replacement
 - $h_i \leftarrow \text{learn}(T_i)$ *[Decision Tree, Naive Bayes, ...]*
- Now combine the h_i together with uniform voting ($w_i = 1/K$ for all i)

Bagging Example




decision tree learning algorithm; very similar to version in earlier slides

CART decision boundary




100 bagged trees



shades of blue/red indicate strength of vote for particular classification

Regression results

Squared error loss



Fighting the bias-variance tradeoff

- **Simple (a.k.a. weak) learners are good**
 - e.g., naïve Bayes, logistic regression, decision stumps (or shallow decision trees)
 - Low variance, don't usually overfit
- **Simple (a.k.a. weak) learners are bad**
 - High bias, can't solve hard learning problems
- **Can we make weak learners always good???**
 - **No!!!**
 - **But often yes...**

Boosting

[Schapire, 1989]

- **Idea:** given a weak learner, run it multiple times on (reweighted) training data, then let learned classifiers vote
- **On each iteration t :**
 - weight each training example by how incorrectly it was classified
 - Learn a hypothesis – h_t
 - A strength for this hypothesis – α_t

- **Final classifier:**

$$h(x) = \text{sign} \left(\sum_i \alpha_i h_i(x) \right)$$

- **Practically useful**
- **Theoretically interesting**

Boosting

[Schapire, 1989]


- **Idea:** given a weak learner, run it multiple times on (reweighted) training data, then let learned classifiers vote
- **On each iteration t :**
 - weight each training example by how incorrectly it was classified
 - Learn a hypothesis – h_t
 - A strength for this hypothesis – α_t

Neural network with one hidden layer?

- **Final classifier:**

$$h(x) = \text{sign} \left(\sum_i \alpha_i h_i(x) \right)$$

- **Practically useful**
- **Theoretically interesting**




time = 0

blue/red = class

size of dot = weight


weak learner =
Decision stub:
horizontal or vertical

First, generate a data-set by clicking on the left and right buttons in the main window of the applet. Then, press "split" to split the data into training and test sets




First, generate a data-set by clicking on the left and right buttons in the main window of the applet. Then, press "split" to split the data into training and test sets.

Applet adaboost started




time = 2

First, generate a data-set by clicking on the left and right buttons in the main window of the applet. Then, press "split" to split the data into training and test sets




time = 3

First, generate a data-set by clicking on the left and right buttons in the main window of the applet. Then, press "split" to split the data into training and test sets



time = 100

First, generate a data-set by clicking on the left and right buttons in the main window of the applet. Then, press "split" to split the data into training and test sets



overfitting

First, generate a data-set by clicking on the left and right buttons in the main window of the applet. Then, press "split" to split the data into training and test sets.

Learning from weighted data

- **Consider a weighted dataset**

- $D(i)$ – weight of i th training example (\mathbf{x}^i, y^i)
- Interpretations:
 - i th training example counts as if it occurred $D(i)$ times
 - If I were to “resample” data, I would get more samples of “heavier” data points

- **Now, always do weighted calculations:**

- e.g., MLE for Naïve Bayes, redefine $Count(Y=y)$ to be **weighted** count:

$$Count(Y = y) = \sum_{j=1}^n D(j) \delta(Y^j = y)$$

- setting $D(j)=1$ (or any constant value!), for all j , will recreate unweighted case

Given: $(x^1, y^1), \dots, (x^m, y^m)$ where $x^i \in \mathbb{R}^n, y^i \in \{-1, +1\}$

Initialize: $D_1(i) = 1/m$, for $i = 1, \dots, m$

How? Many possibilities. Will see one shortly!

For $t=1 \dots T$:

- Train base classifier $h_t(x)$ using D_t
- Choose α_t
- Update, for $i=1 \dots m$:

Why? Reweight the data: examples i that are misclassified will have higher weights!

$$D_{t+1}(i) \propto D_t(i) \exp(-\alpha_t y^i h_t(x^i))$$

with normalization constant:

$$\sum_{i=1}^m D_t(i) \exp(-\alpha_t y^i h_t(x^i))$$

- $y^i h_t(x^i) > 0 \rightarrow h_i$ correct
- $y^i h_t(x^i) < 0 \rightarrow h_i$ wrong
- h_i correct, $\alpha_t > 0 \rightarrow D_{t+1}(i) < D_t(i)$
- h_i wrong, $\alpha_t > 0 \rightarrow D_{t+1}(i) > D_t(i)$

Output final classifier:

$$H(x) = \text{sign} \left(\sum_{i=1}^T \alpha_t h_t(x) \right)$$

Final Result: linear sum of “base” or “weak” classifier outputs.

Given: $(x^1, y^1), \dots, (x^m, y^m)$ where $\epsilon_t = \sum_{i=1}^m D_t(i) \delta(h_t(x^i) \neq y^i)$

Initialize: $D_1(i) = 1/m$, for $i = 1, \dots, m$


For $t=1 \dots T$:

- Train base classifier $h_t(x)$ using D_t
- Choose α_t
- Update, for $i=1 \dots m$:

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$D_{t+1}(i) \propto D_t(i) \exp(-\alpha_t y^i h_t(x^i))$$

- ϵ_t : error of h_t , weighted by D_t
 - $0 \leq \epsilon_t \leq 1$
- α_t :
 - No errors: $\epsilon_t=0 \rightarrow \alpha_t=\infty$
 - All errors: $\epsilon_t=1 \rightarrow \alpha_t=-\infty$
 - Random: $\epsilon_t=0.5 \rightarrow \alpha_t=0$




What α_t to choose for hypothesis h_t ?

[Schapire, 1989]

Idea: choose α_t to minimize a bound on training error!

$$\frac{1}{m} \sum_{i=1}^m \delta(H(x^i) \neq y^i) \leq \sum_{i=1}^m D_t(i) \exp(-y^i f(x^i))$$


Where $f(x) = \sum_t \alpha_t h_t(x)$; $H(x) = \text{sign}(f(x))$



What α_t to choose for hypothesis h_t ?

[Schapire, 1989]

Idea: choose α_t to minimize a bound on training error!

$$\frac{1}{m} \sum_{i=1}^m \delta(H(x^i) \neq y^i) \leq \sum_{i=1}^m D_t(i) \exp(-y^i f(x^i)) = \prod_t Z_t$$


Where

$$f(x) = \sum_t \alpha_t h_t(x); H(x) = \text{sign}(f(x))$$

And

$$Z_t = \sum_{i=1}^m D_t(i) \exp(-\alpha_t y^i h_t(x^i))$$

This equality isn't obvious! Can be shown with algebra (telescoping sums)!

If we minimize $\prod_t Z_t$, we minimize our training error!!!

- We can tighten this bound greedily, by choosing α_t and h_t on each iteration to minimize Z_t .
- h_t is estimated as a black box, but can we solve for α_t ?

Summary: choose α_t to minimize *error bound*

[Schapire, 1989]

We can squeeze this bound by choosing α_t on each iteration to minimize Z_t .

$$Z_t = \sum_{i=1}^m D_t(i) \exp(-\alpha_t y^i h_t(x^i))$$
$$\epsilon_t = \sum_{i=1}^m D_t(i) \delta(h_t(x^i) \neq y^i)$$

For boolean Y: differentiate, set equal to 0, **there is a closed form solution!** [Freund & Schapire '97]:

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$

Initialize: $D_1(i) = 1/m$, for $i = 1, \dots, m$

For $t=1 \dots T$:

- Train base classifier $h_t(x)$ using D_t
- Choose $\alpha_t = \sum_{i=1}^m D_t(i) \delta(h_t(x^i) \neq y^i)$

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$
- Update, for $i=1 \dots m$:

$$D_{t+1}(i) \propto D_t(i) \exp(-\alpha_t y^i h_t(x^i))$$

Output final classifier:

$$H(x) = \text{sign} \left(\sum_{i=1}^T \alpha_i h_i(x) \right)$$

x_1	y
-1	1
0	-1
1	1



$$H(x) = \text{sign}(0.35 \times h_1(x))$$

- $h_1(x) = +1$ if $x_1 > 0.5$, -1 otherwise

Use decision stubs as base classifier

Initial:

- $D_1 = [D_1(1), D_1(2), D_1(3)] = [.33, .33, .33]$

$t=1$:

- Train stub [work omitted, breaking ties randomly]
 - $h_1(x) = +1$ if $x_1 > 0.5$, -1 otherwise
- $\epsilon_1 = \sum_i D_1(i) \delta(h_1(x^i) \neq y^i)$
 $= 0.33 \times 1 + 0.33 \times 0 + 0.33 \times 0 = 0.33$
- $\alpha_1 = (1/2) \ln((1 - \epsilon_1)/\epsilon_1) = 0.5 \times \ln(2) = 0.35$
- $D_2(1) \propto D_1(1) \times \exp(-\alpha_1 y^1 h_1(x^1))$
 $= 0.33 \times \exp(-0.35 \times 1 \times -1) = 0.33 \times \exp(0.35) = 0.46$
- $D_2(2) \propto D_1(2) \times \exp(-\alpha_1 y^2 h_1(x^2))$
 $= 0.33 \times \exp(-0.35 \times -1 \times -1) = 0.33 \times \exp(-0.35) = 0.23$
- $D_2(3) \propto D_1(3) \times \exp(-\alpha_1 y^3 h_1(x^3))$
 $= 0.33 \times \exp(-0.35 \times 1 \times 1) = 0.33 \times \exp(-0.35) = 0.23$
- $D_2 = [D_2(1), D_2(2), D_2(3)] = [0.46, 0.23, 0.23]$

$t=2$

- Continues on next slide!

Initialize: $D_1(i) = 1/m$, for $i = 1, \dots, m$

For $t=1 \dots T$:

- Train base classifier $h_t(x)$ using D_t
- Choose $\alpha_t = \sum_{i=1}^m D_t(i) \delta(h_t(x^i) \neq y^i)$

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$
- Update, for $i=1 \dots m$:

$$D_{t+1}(i) \propto D_t(i) \exp(-\alpha_t y^i h_t(x^i))$$

Output final classifier:

$$H(x) = \text{sign} \left(\sum_{i=1}^T \alpha_i h_i(x) \right)$$

x_1	y
-1	1
0	-1
1	1



$$H(x) = \text{sign}(0.35 \times h_1(x) + 0.55 \times h_2(x))$$

- $h_1(x) = +1$ if $x_1 > 0.5$, -1 otherwise
- $h_2(x) = +1$ if $x_1 < 1.5$, -1 otherwise

- $D_2 = [D_2(1), D_2(2), D_2(3)] = [0.5, 0.25, 0.25]$

$t=2$:

- Train stub [work omitted; different stub because of new data weights D ; breaking ties opportunistically (will discuss at end)]
 - $h_2(x) = +1$ if $x_1 < 1.5$, -1 otherwise
- $\epsilon_2 = \sum_i D_2(i) \delta(h_2(x^i) \neq y^i)$
 $= 0.5 \times 0 + 0.25 \times 1 + 0.25 \times 0 = 0.25$
- $\alpha_2 = (1/2) \ln((1 - \epsilon_2)/\epsilon_2) = 0.5 \times \ln(3) = 0.55$
- $D_2(1) \propto D_1(1) \times \exp(-\alpha_2 y^1 h_2(x^1))$
 $= 0.5 \times \exp(-0.55 \times 1 \times 1) = 0.5 \times \exp(-0.55) = 0.29$
- $D_2(2) \propto D_1(2) \times \exp(-\alpha_2 y^2 h_2(x^2))$
 $= 0.25 \times \exp(-0.55 \times -1 \times 1) = 0.25 \times \exp(0.55) = 0.43$
- $D_2(3) \propto D_1(3) \times \exp(-\alpha_2 y^3 h_2(x^3))$
 $= 0.25 \times \exp(-0.55 \times 1 \times 1) = 0.25 \times \exp(-0.55) = 0.14$
- $D_3 = [D_3(1), D_3(2), D_3(3)] = [0.33, 0.5, 0.17]$

$t=3$

- Continues on next slide!

Initialize: $D_1(i) = 1/m$, for $i = 1, \dots, m$

For $t=1 \dots T$:

- Train base classifier $h_t(x)$ using D_t
- Choose $\alpha_t = \sum_{i=1}^m D_t(i) \delta(h_t(x^i) \neq y^i)$

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$
- Update, for $i=1 \dots m$:

$$D_{t+1}(i) \propto D_t(i) \exp(-\alpha_t y^i h_t(x^i))$$

Output final classifier:

$$H(x) = \text{sign} \left(\sum_{i=1}^T \alpha_i h_i(x) \right)$$

x_1	y
-1	1
0	-1
1	1



$$H(x) = \text{sign}(0.35 \times h_1(x) + 0.55 \times h_2(x) + 0.79 \times h_3(x))$$

- $h_1(x) = +1$ if $x_1 > 0.5$, -1 otherwise
- $h_2(x) = +1$ if $x_1 < 1.5$, -1 otherwise
- $h_3(x) = +1$ if $x_1 < -0.5$, -1 otherwise

$$D_3 = [D_3(1), D_3(2), D_3(3)] = [0.33, 0.5, 0.17]$$

$t=3$:

- Train stub [work omitted; different stub because of new data weights D ; breaking ties opportunistically (will discuss at end)]
 - $h_3(x) = +1$ if $x_1 < -0.5$, -1 otherwise
- $\epsilon_3 = \sum_i D_3(i) \delta(h_3(x^i) \neq y^i)$

$$= 0.33 \times 0 + 0.5 \times 0 + 0.17 \times 1 = 0.17$$
- $\alpha_3 = (1/2) \ln((1 - \epsilon_3)/\epsilon_3) = 0.5 \times \ln(4.88) = 0.79$
- Stop!!! How did we know to stop?

Strong, weak classifiers

- If each classifier is (at least slightly) better than random: $\epsilon_t < 0.5$


- Another bound on error:

$$\frac{1}{m} \sum_{i=1}^m \delta(H(x^i) \neq y^i) \leq \prod_t Z_t \leq \exp \left(-2 \sum_{t=1}^T (1/2 - \epsilon_t)^2 \right)$$

- What does this imply about the training error?
 - Will reach zero!
 - Will get there exponentially fast!
- Is it hard to achieve better than random training error?

Boosting results – Digit recognition

[Schapire, 1989]





- **Boosting:**
 - Seems to be robust to overfitting
 - Test error can decrease even after training error is zero!!!

Boosting: Experimental Results

[Freund & Schapire, 1996]

Comparison of C4.5, Boosting C4.5, Boosting decision stumps (depth 1 trees), 27 benchmark datasets




Boosting and Logistic Regression

Logistic regression equivalent
to minimizing log loss:

Boosting minimizes similar
loss function:

$$\ln(1 + \exp(-y^i f(x^i)))$$

$$\exp(-y^i f(x^i))$$



Both smooth approximations of 0/1 loss!

Logistic regression and Boosting

Logistic regression:

- Minimize loss fn

$$\sum_{i=1}^m \ln(1 + \exp(-y^i f(x^i)))$$

- Define

$$f(x) = \sum_j w_j x_j$$

where each feature x_j is predefined

- Jointly optimize parameters w_0, w_1, \dots, w_n via gradient ascent.

Boosting:

- Minimize loss fn

$$\sum_{i=1}^m \exp(-y^i f(x^i))$$

- Define

$$f(x) = \sum_t \alpha_t h_t(x)$$

where $h_t(x)$ learned to fit data

- Weights α_j learned incrementally (new one for each training pass)

What you need to know about Boosting

- Combine weak classifiers to get very strong classifier
 - Weak classifier – slightly better than random on training data
 - Resulting very strong classifier – can get zero training error
- AdaBoost algorithm
- Boosting v. Logistic Regression
 - Both linear model, boosting “learns” features
 - Similar loss functions
 - Single optimization (LR) v. Incrementally improving classification (B)
- Most popular application of Boosting:
 - Boosted decision stumps!
 - Very simple to implement, very effective classifier